

Merge and Code Review

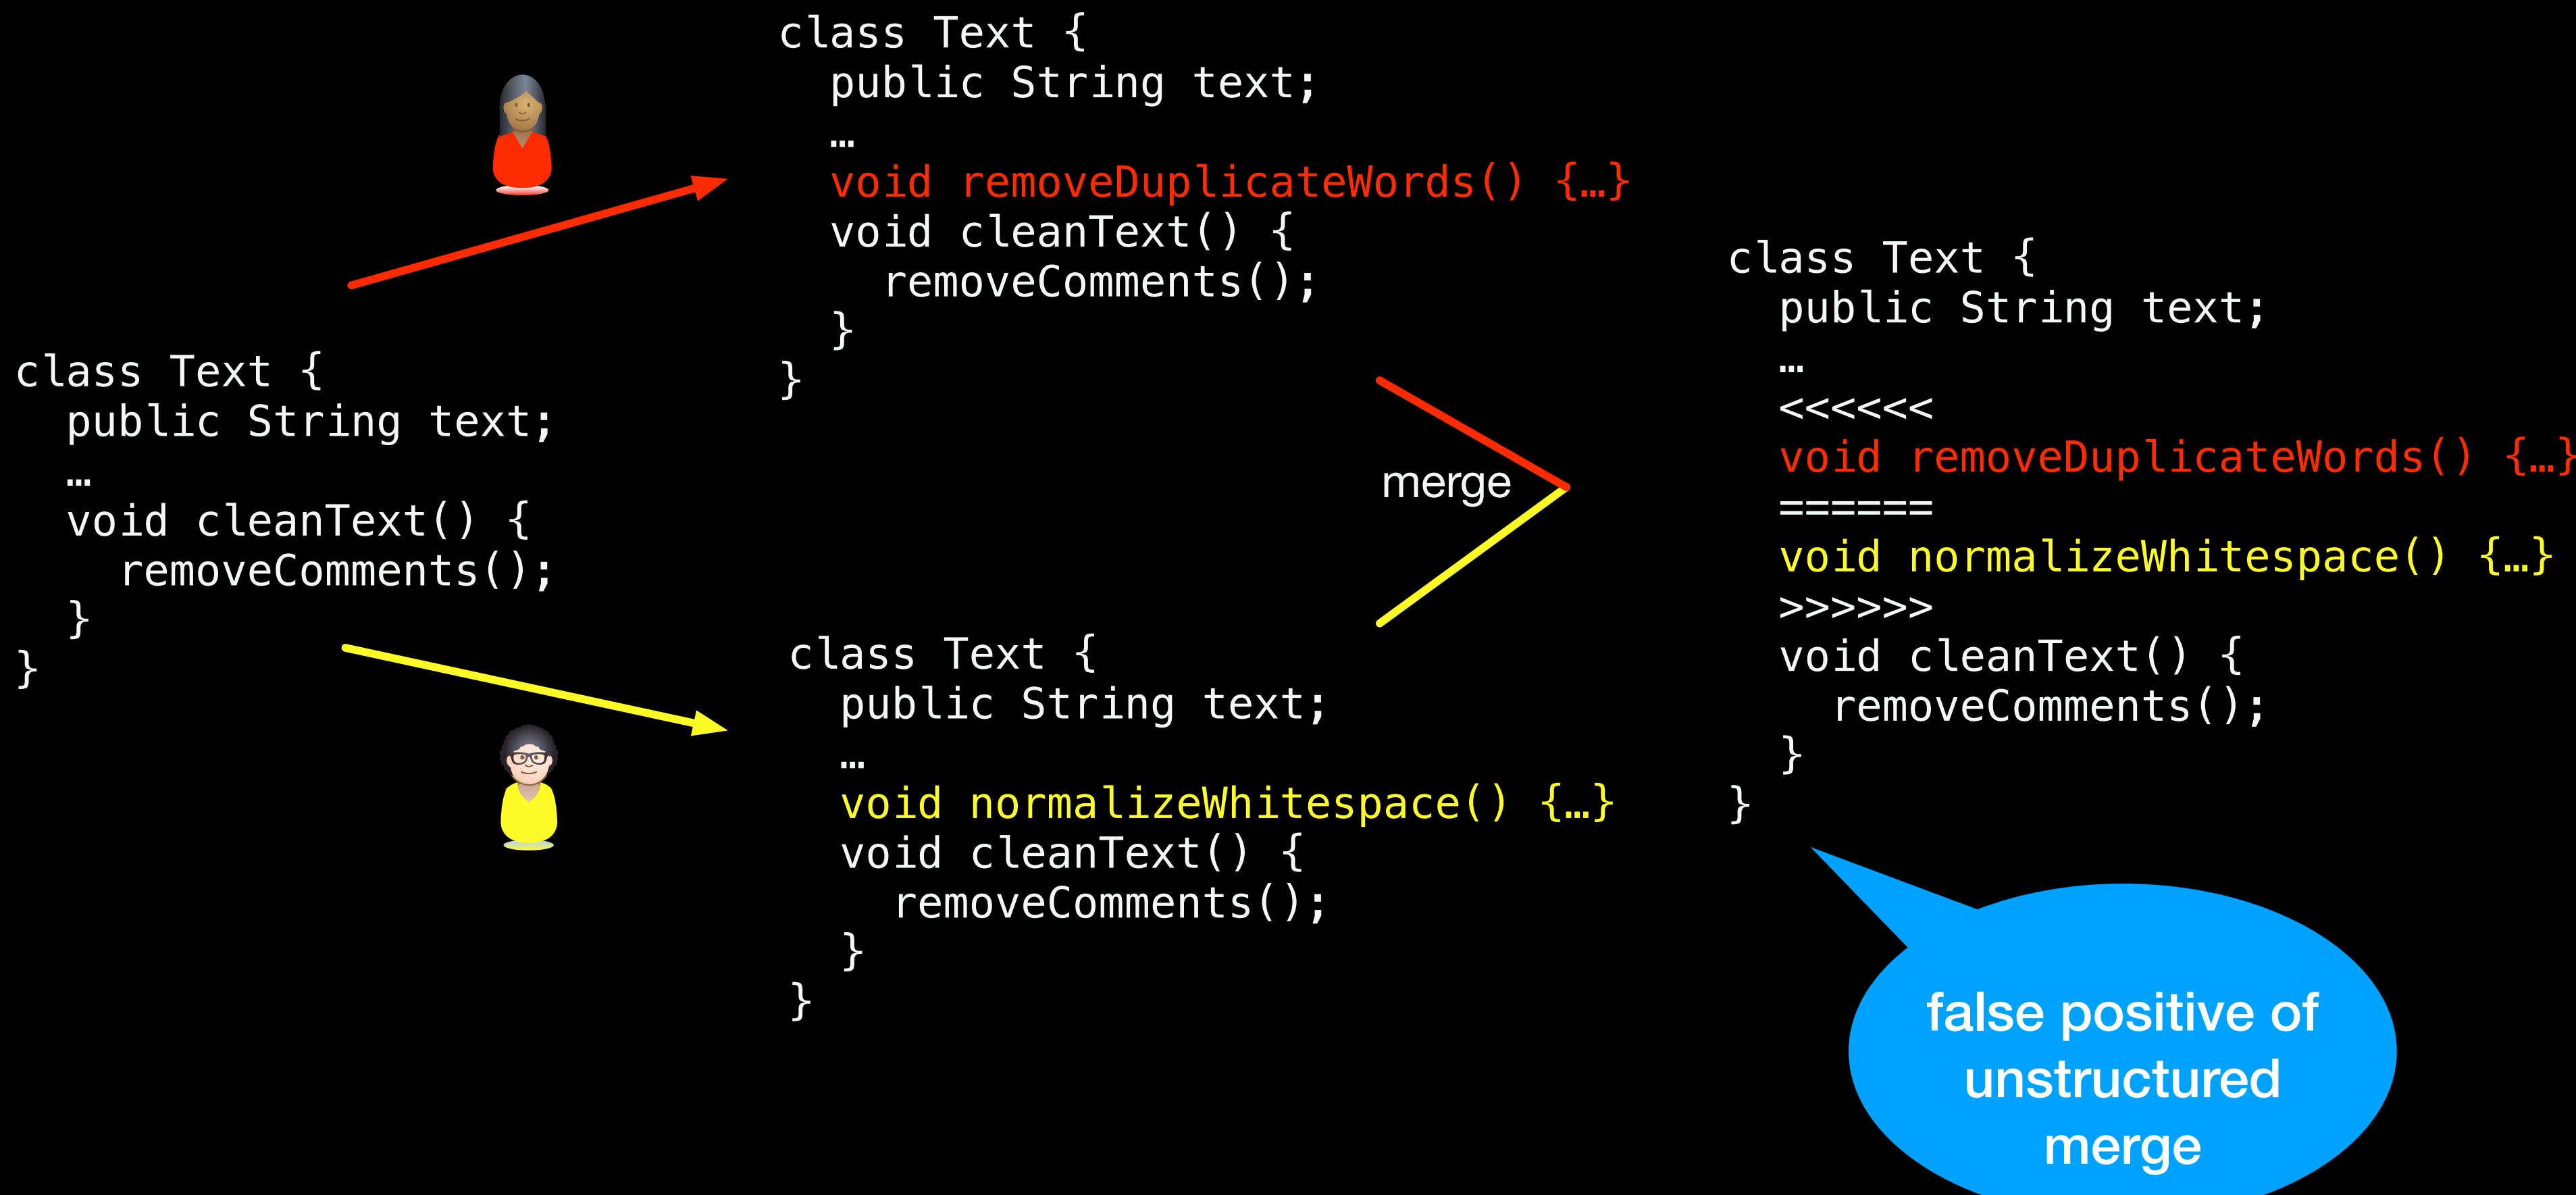
Paulo Borba
Informatics Center
Federal University of Pernambuco

pauloborba.cin.ufpe.br

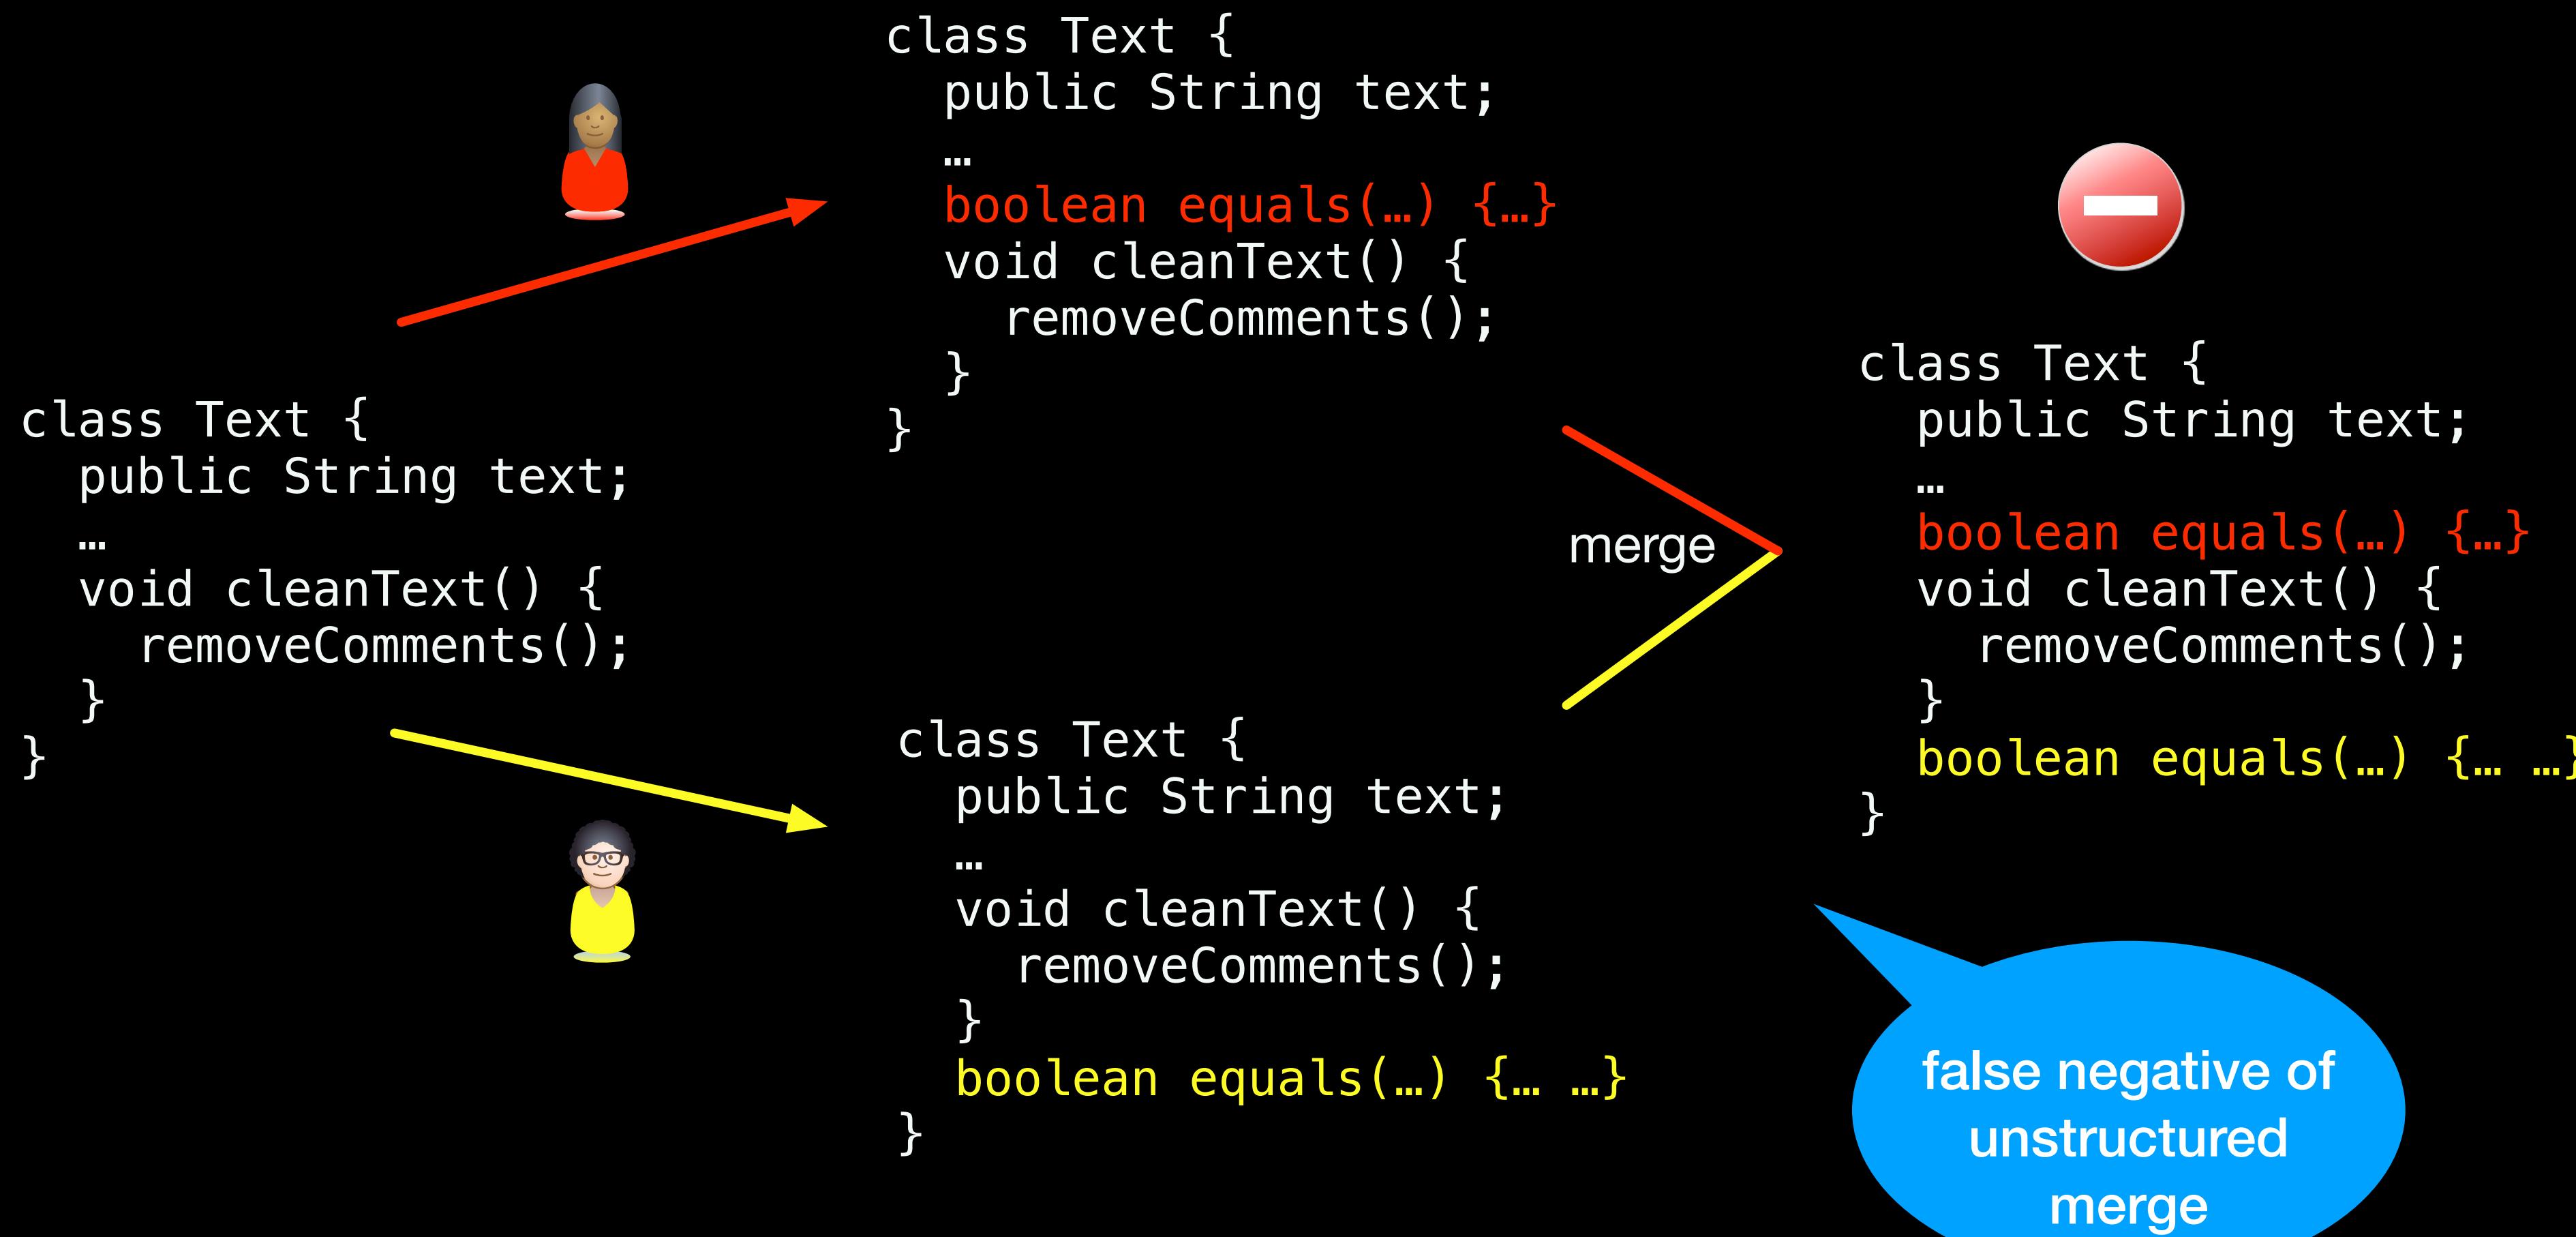
Semistructured merge

Trying to avoid unstructured merge
false positives and negatives

Developers waste time by manually resolving conflicts that could be automatically solved



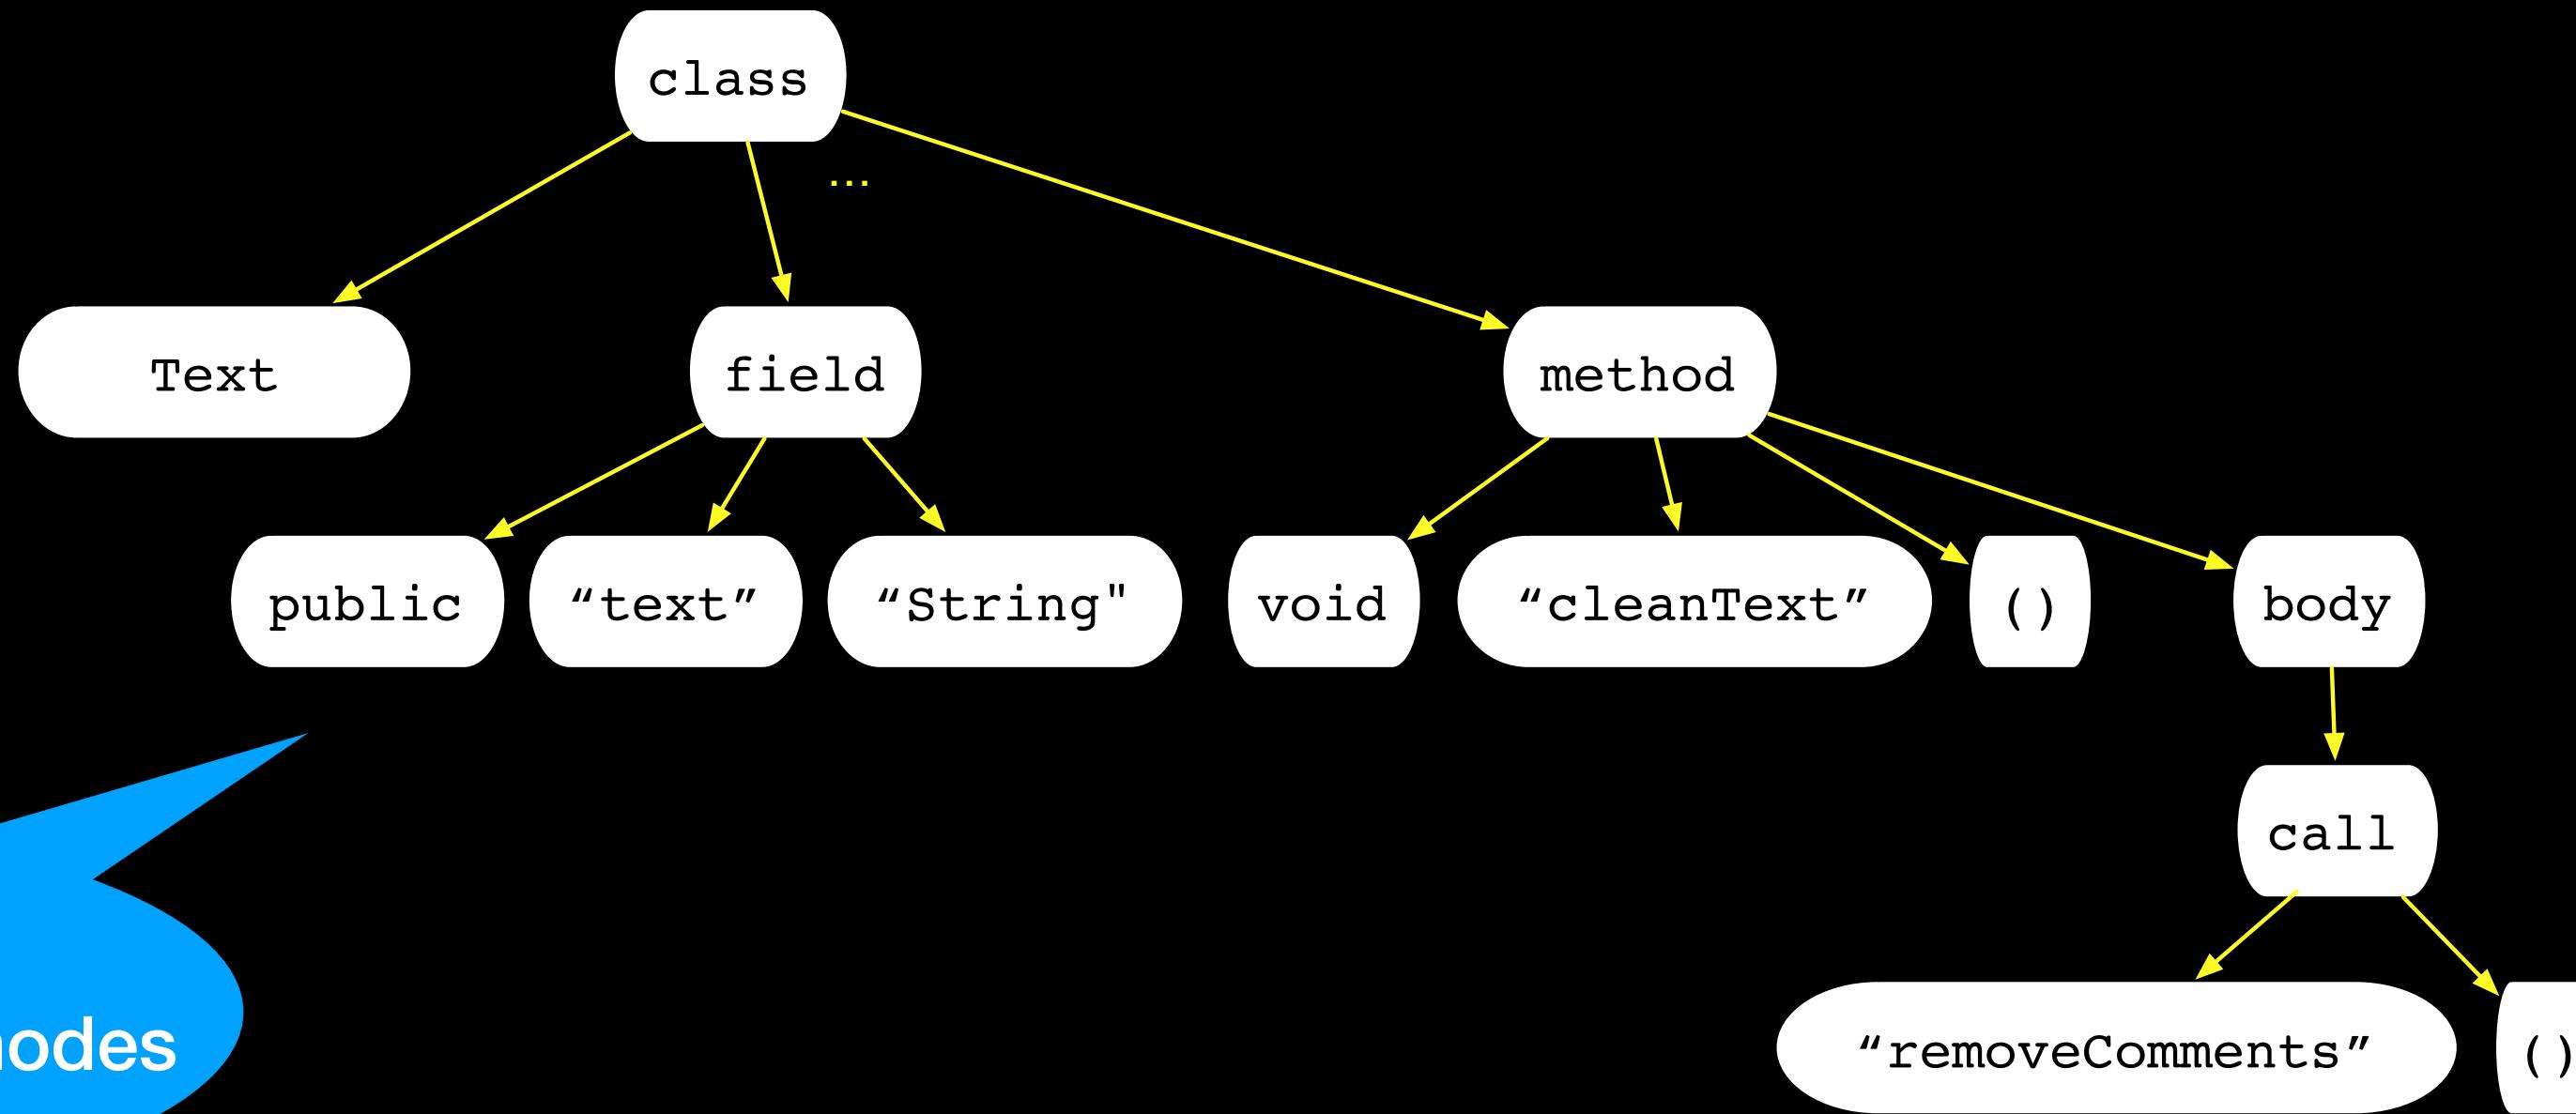
Current merge tools might integrate conflicting changes without warning developers



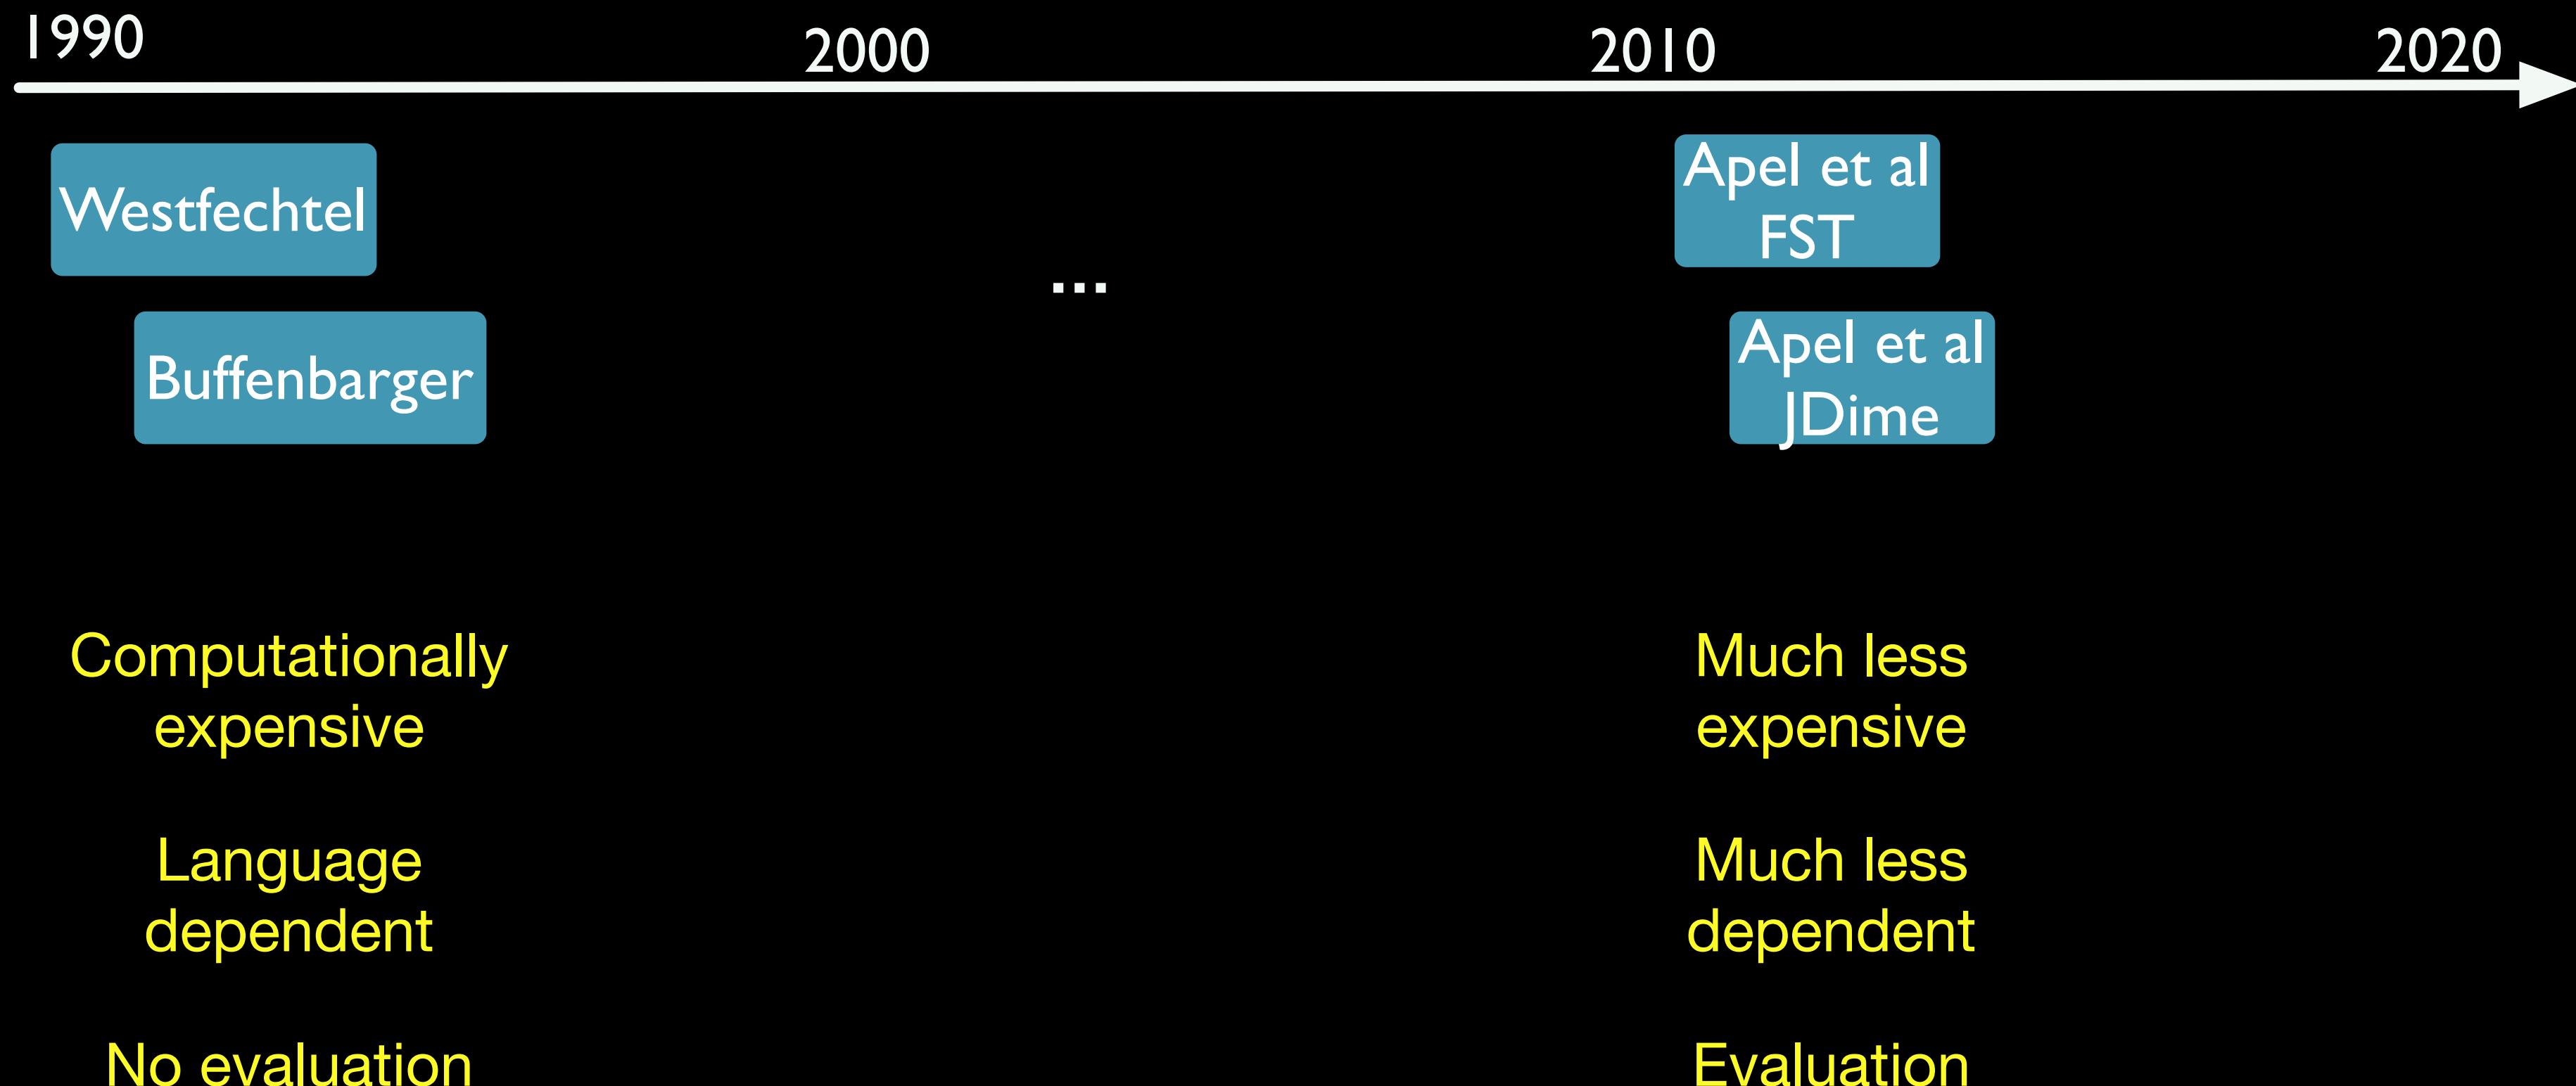
**Structured
and
semistructured
merge tools to the rescue**

Structured merge works with ASTs, not a sequence of lines

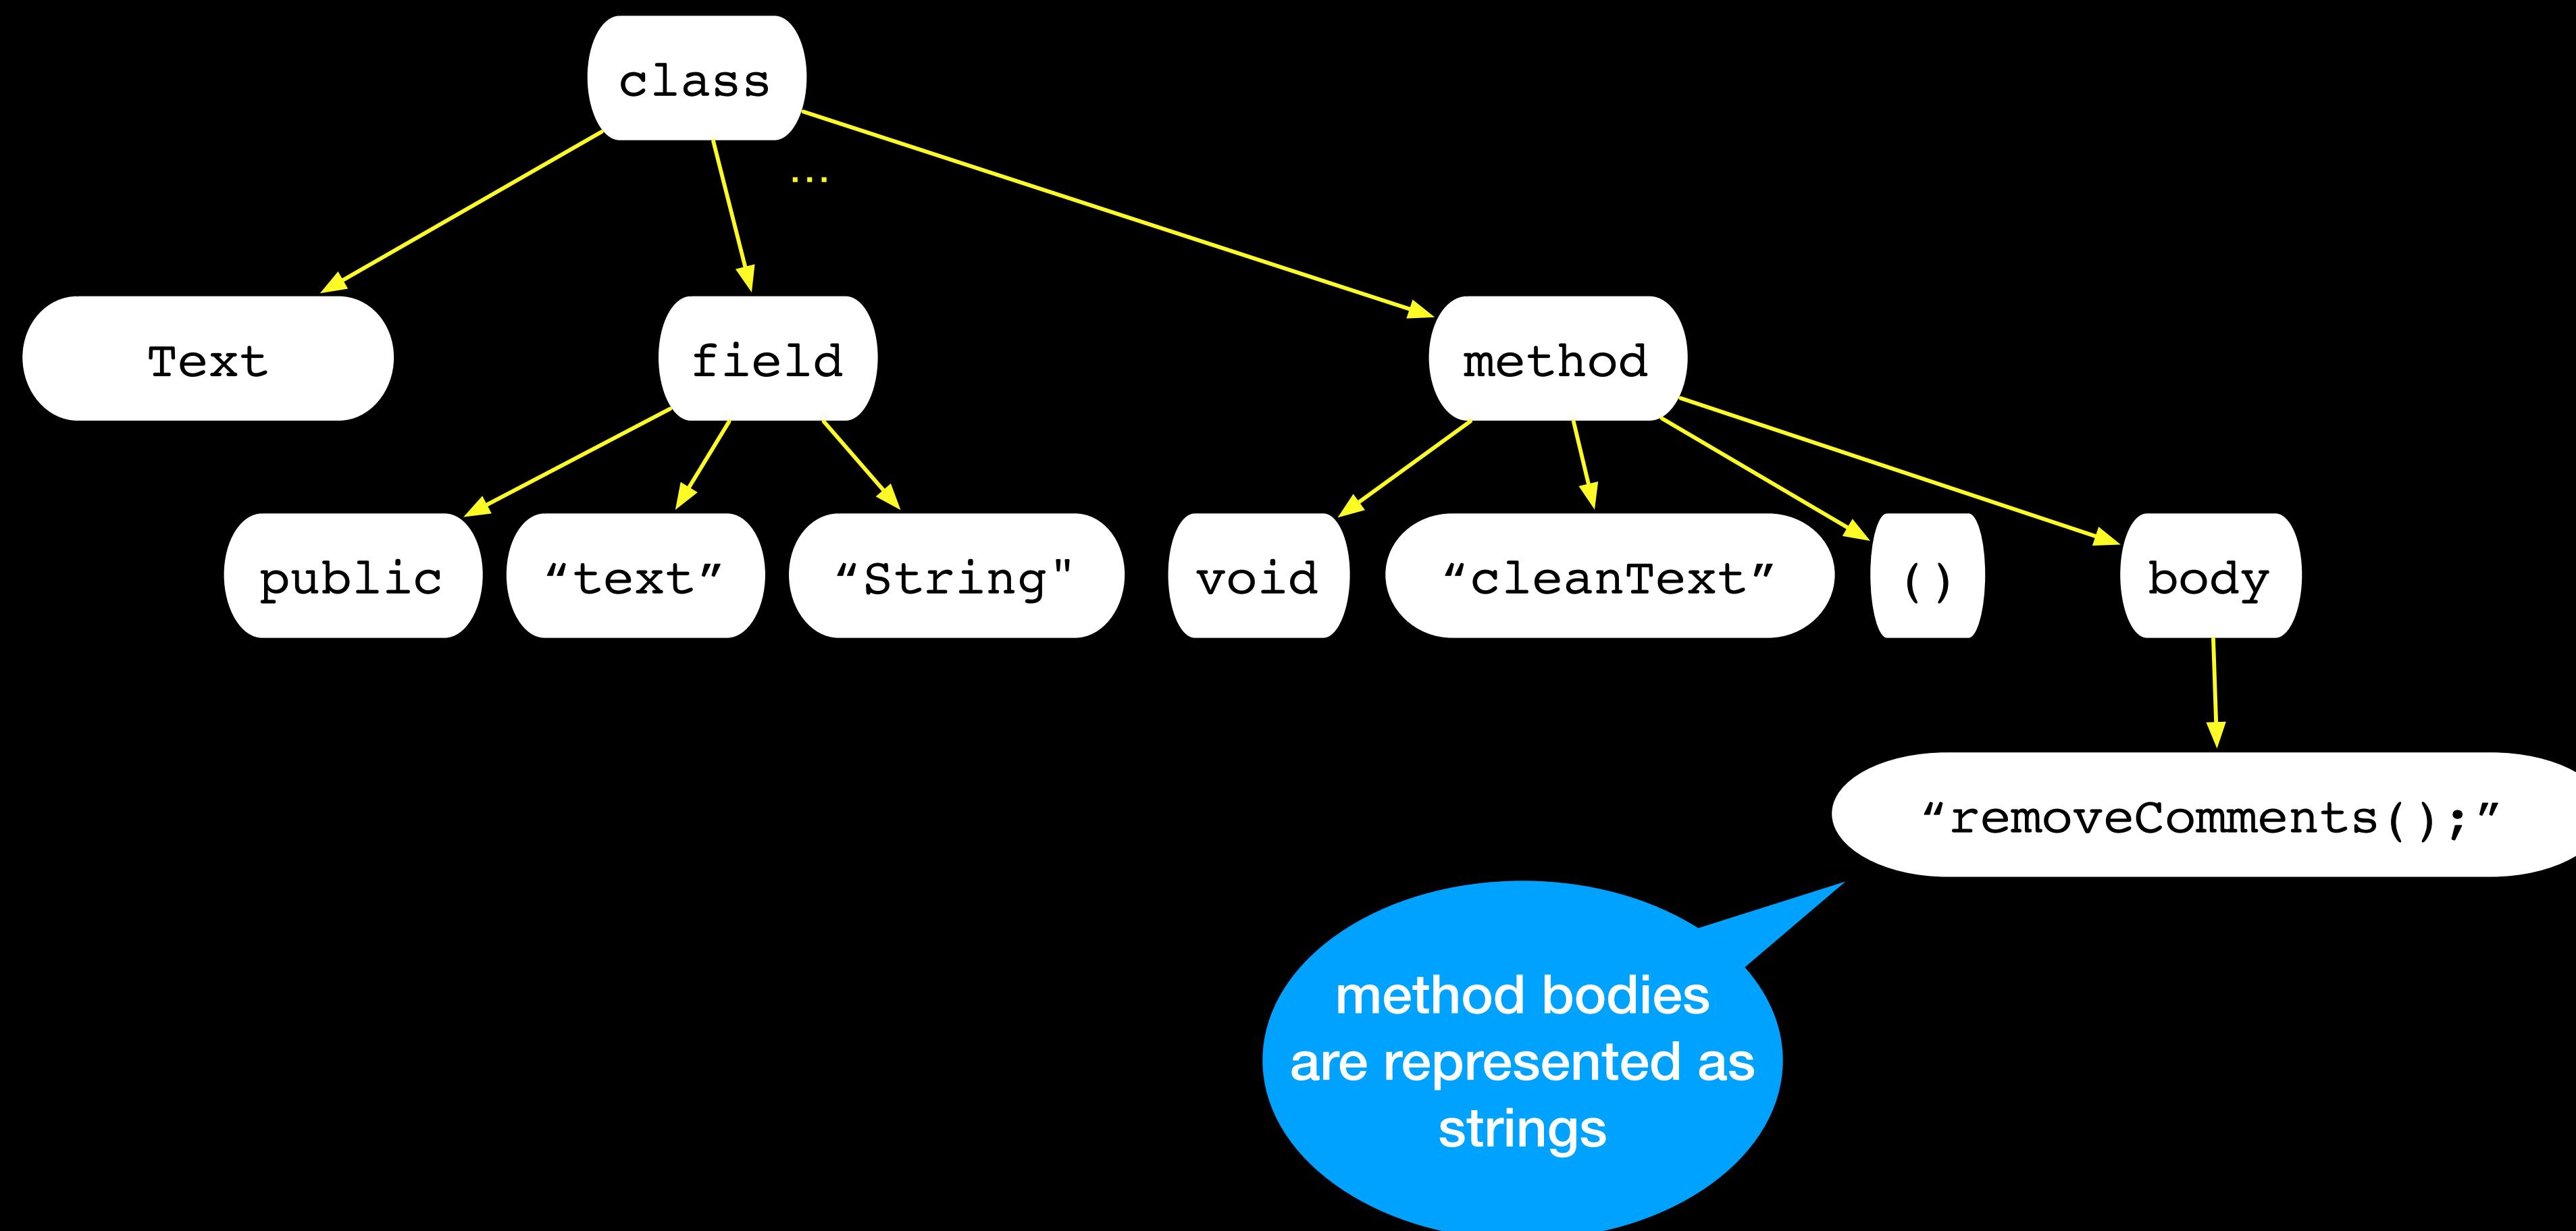
```
class Text {  
    public String text;  
    ...  
    void cleanText() {  
        removeComments();  
    }  
}
```



Structured merge tools timeline

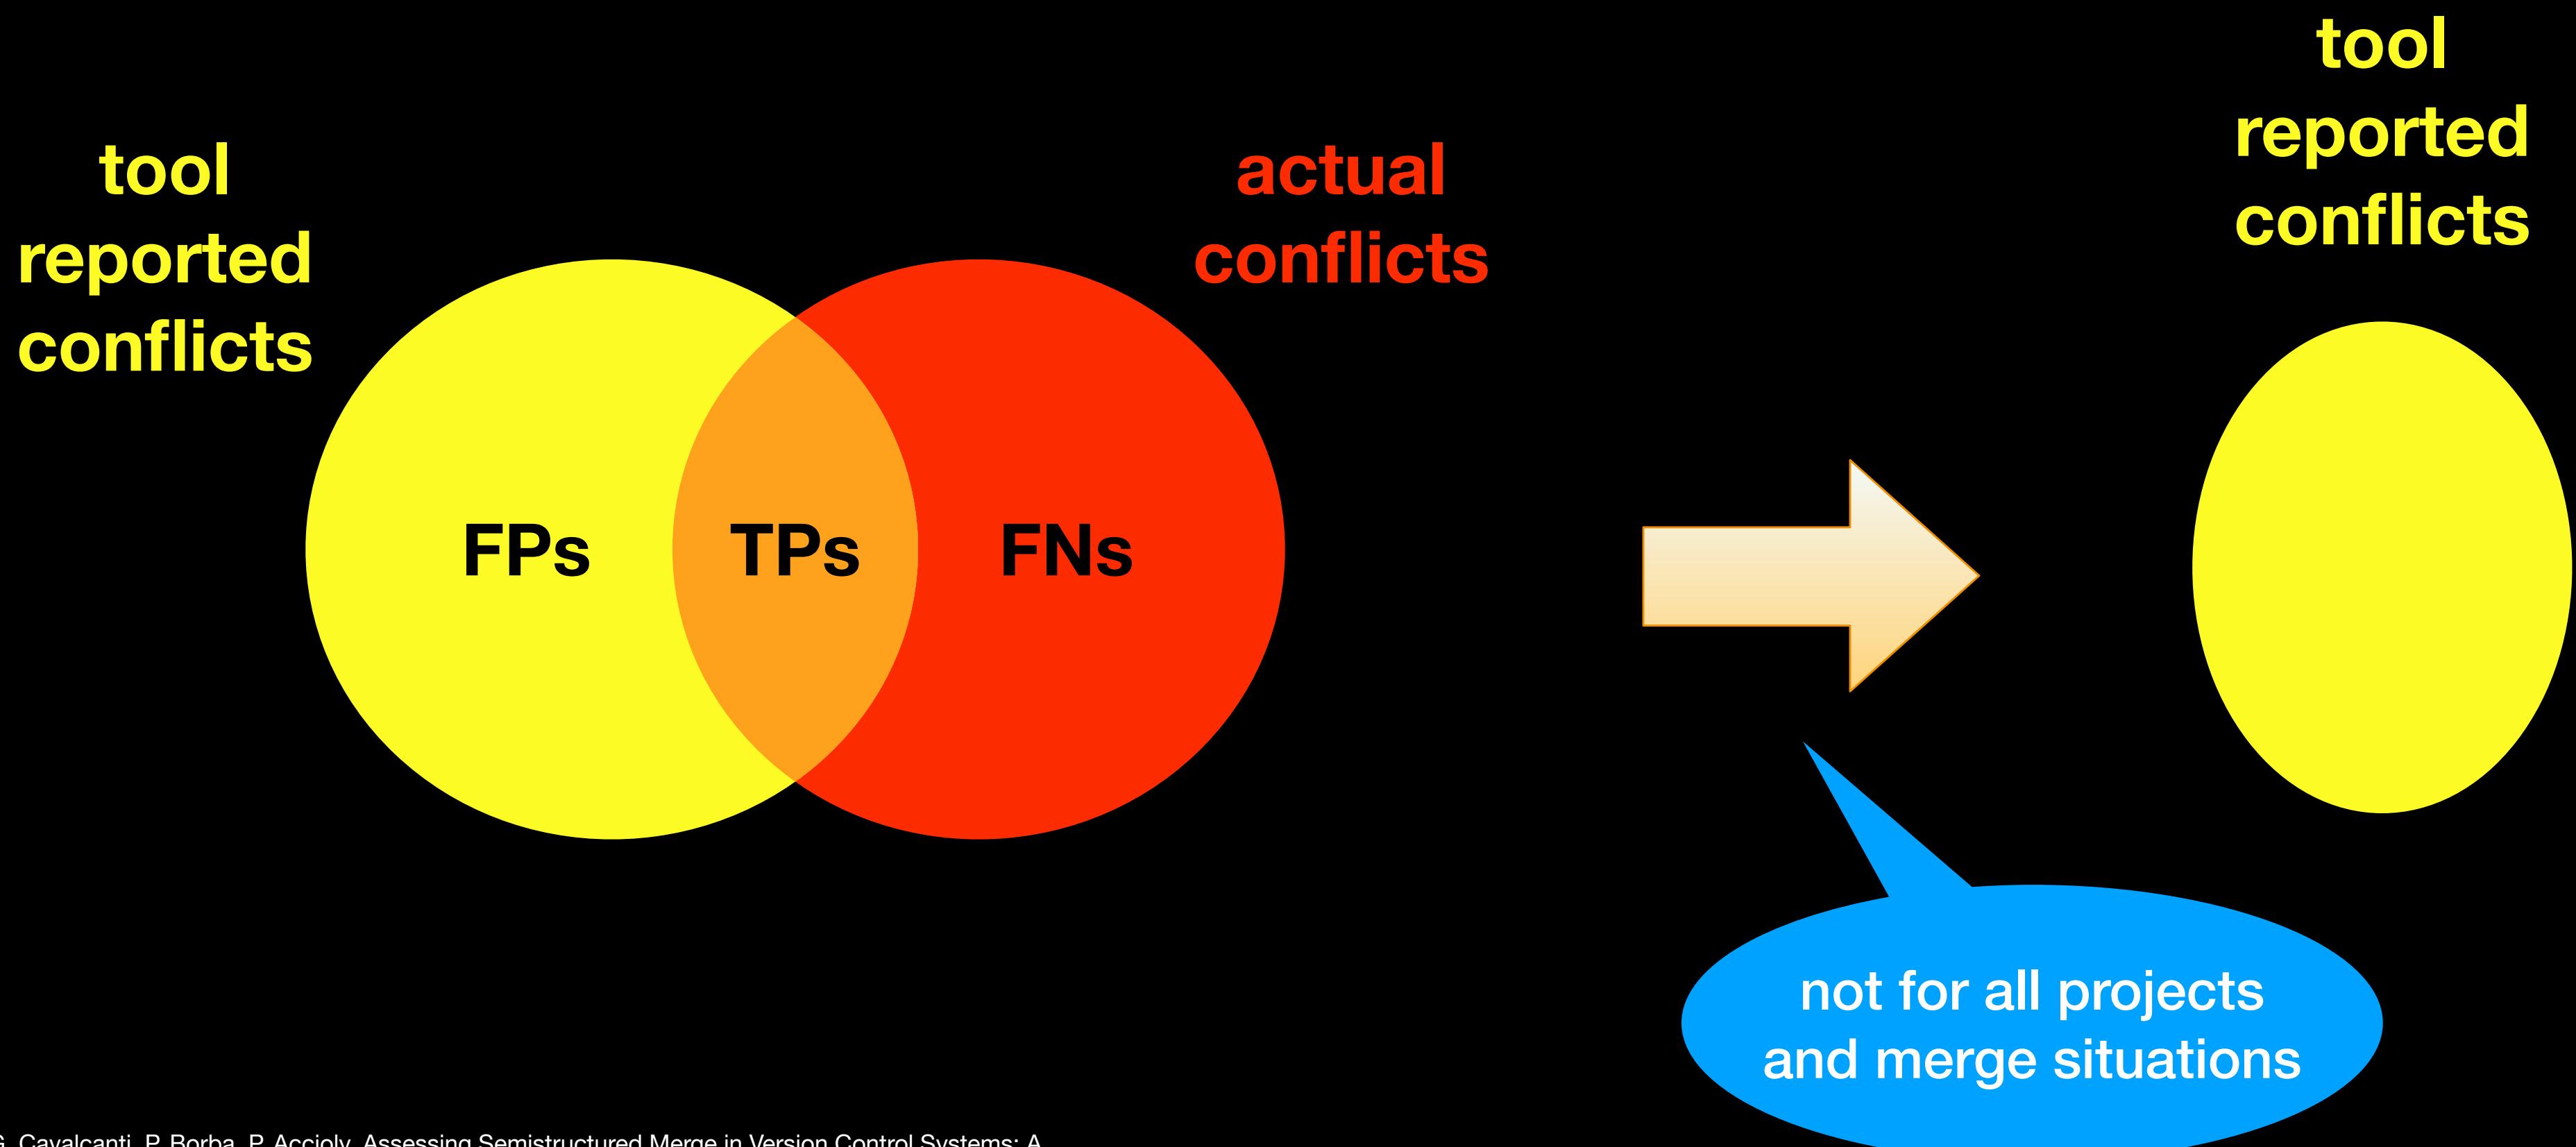


Semistructured merge works with partial ASTs

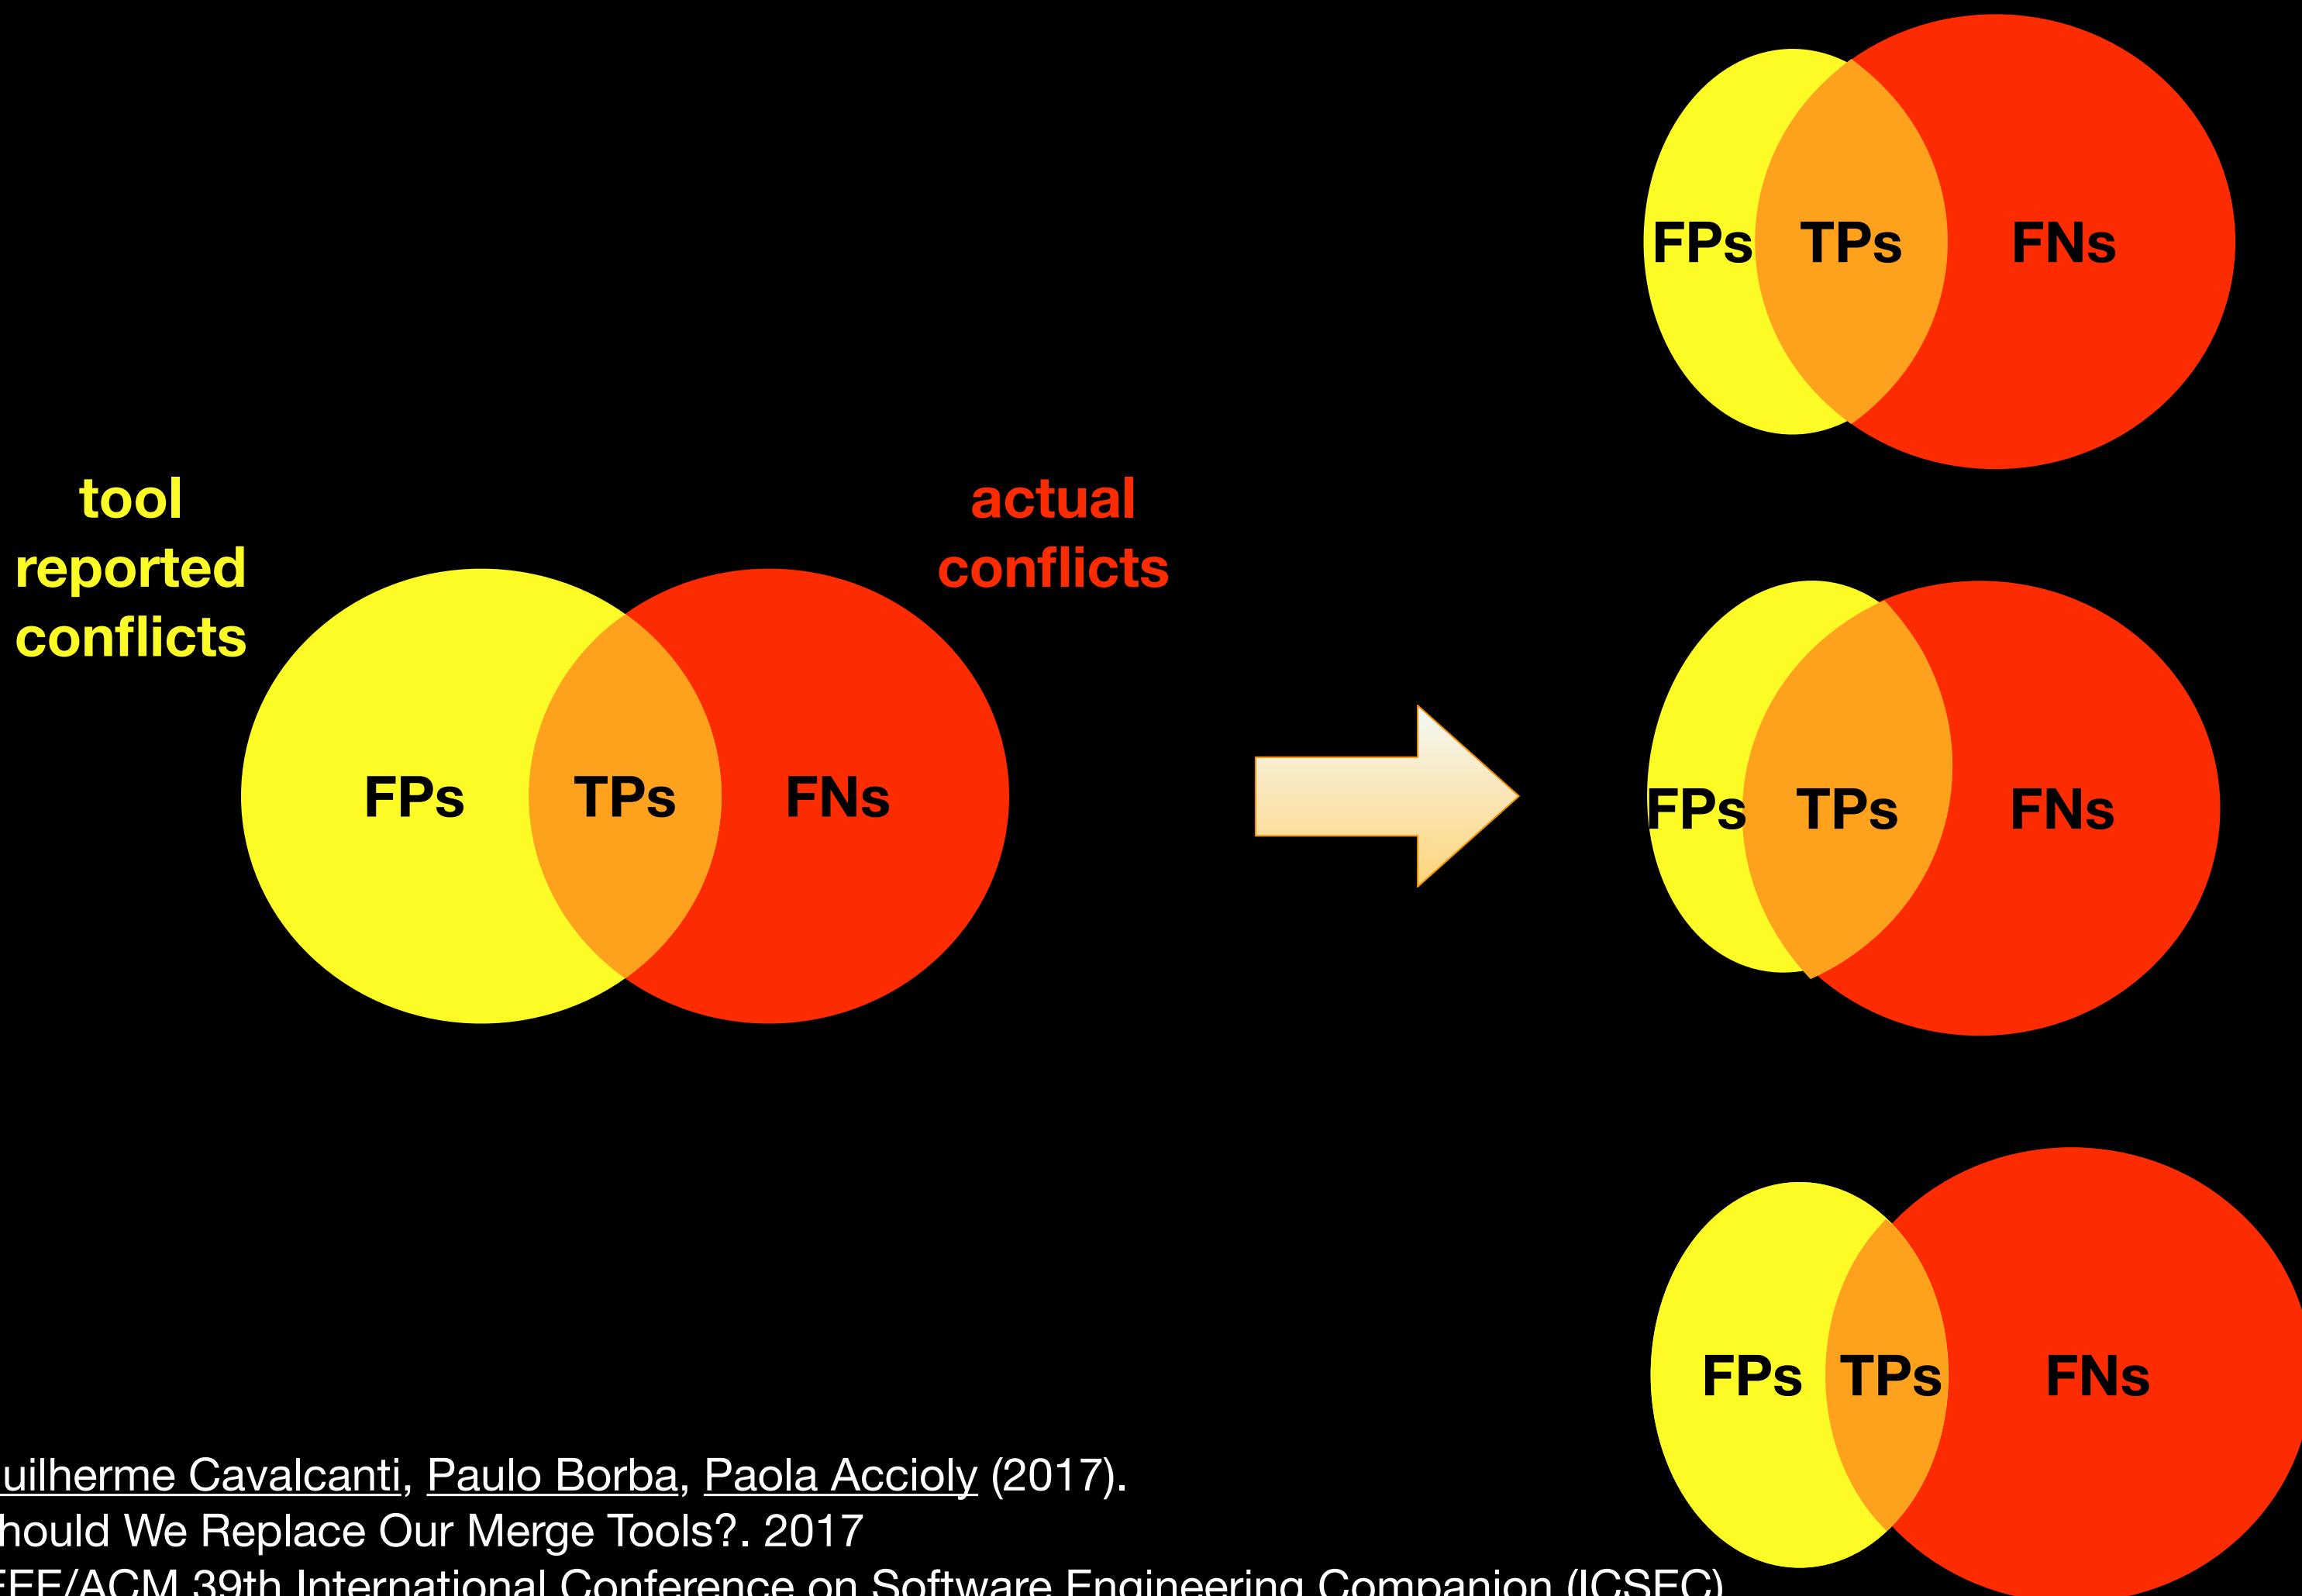


Evidence only about the reduction of reported conflicts

(others: $34\% \pm 21\%$, 39%) (here: $62\% \pm 24\%$, $71\% \pm 30\%$)



Integration effort reduction? Integration correctness impact?



A more comprehensive evaluation is not fully favourable for semistructured merge

(34,030 merges from 50 open source projects)

FSTMerge reports **less false positives** (34%)
than unstructured merge tools.

Easier to analyze and resolve.

No evidence that FSTMerge leads to **fewer
false negatives**.

Harder to detect and resolve.

not uniformly
across projects



Evaluating and Improving Semistructured Merge

GUILHERME CAVALCANTI, Federal University of Pernambuco, Brazil
PAULO BORBA, Federal University of Pernambuco, Brazil
PAOLA ACCIOLY, Federal University of Pernambuco, Brazil

While unstructured merge tools rely only on textual analysis to detect and resolve conflicts, semistructured merge tools go further by partially exploiting the syntactic structure and semantics of the involved artifacts. Previous studies compare these merge approaches with respect to the number of reported conflicts, showing, for most projects and merge situations, reduction in favor of semistructured merge. However, these studies do not investigate whether this reduction actually leads to integration effort reduction (productivity) without negative impact on the correctness of the merging process (quality). To analyze that, and better understand how merge tools could be improved, in this paper we reproduce more than 30,000 merges from 50 open source projects, identifying conflicts incorrectly reported by one approach but not by the other (false positives), and conflicts correctly reported by one approach but missed by the other (false negatives). Our results and complementary analysis indicate that, in the studied sample, the number of false positives is significantly reduced when using semistructured merge. We also find evidence that its false positives are easier to analyze and resolve than those reported by unstructured merge. However, we find no evidence that semistructured merge leads to fewer false negatives, and we argue that they are harder to detect and resolve than unstructured merge false negatives. Driven by these findings, we implement an improved semistructured merge tool that further combines both approaches to reduce the false positives and false negatives of semistructured merge. We find evidence that the improved tool, when compared to unstructured merge in our sample, reduces the number of reported conflicts by half, has no additional false positives, has at least 8% fewer false negatives, and is not prohibitively slower.

s3m: an improved semistructured merge tool for Java

s3m reports **less conflicts (51%)**

no additional **false positives**

at least 8% **fewer false negatives**

not prohibitively **slower**

(32x, < 1s in 80% of the scenarios, > 5s in only 2%)

But the benefits do not generalize, as strongly, to Javascript (10,345 merges from 50 projects)

s3m reports **less conflicts** (6%)

fewer false positives (87%) without
compromising correctness (1 FN)

commutative and non-commutative
elements at the same syntactic level

Semistructured Merge in JavaScript Systems

Alberto Trindade Tavares, Paulo Borba, Guilherme Cavalcanti, Sérgio Soares
Federal University of Pernambuco
(att, phmb, gic, scbs)@cin.ufpe.br

Abstract—Industry widely uses unstructured merge tools that rely on heuristics to detect conflicts. In this paper, we argue that code contributions. Semistructured merge tools go further by partially exploring the syntactic structure of code artifacts, and, as a consequence, obtaining significant merge accuracy gains for Java-like languages. To understand whether semistructured merge tools can also gain generalization to other kinds of languages, we implement and semistructure tools for JavaScript, and compare them to an unstructured tool. We find that current semistructured merge algorithms and frameworks are not directly applicable for scripting languages like JavaScript. By contrast, we show that our semistructured merge tool, s3m, reports fewer spurious conflicts than unstructured merge, without compromising the correctness of the merging process. The gains, however, are much smaller than the ones observed for Java-like languages. This suggests that semistructured merge advantages might be limited for languages that allow both commutative and non-commutative declarations at the same syntactic level.

Index Terms—Collaborative development, Software merging, Semistructured merge, Version control systems, JavaScript

A. Tavares, P. Borba, G. Cavalcanti, S. Soares. Semistructured Merge in
JavaScript Systems. ASE 2019 https://pauloborba.cin.ufpe.br/publication/2019semistructured_merge_in_javascript_systems/

statement

```
if ((window.onanimationend === undefined) && (window.onwebkitanimationend
    CSS_PREFIX = '-webkit-';
    ANIMATION_PROP = 'WebkitAnimation';
    ANIMATIONEND_EVENT = 'webkitAnimationEnd animationend';
} else {
    ANIMATION_PROP = 'animation';
    ANIMATIONEND_EVENT = 'animationend';
}

var DURATION_KEY = 'Duration';
var PROPERTY_KEY = 'Property';
var DELAY_KEY = 'Delay';
var TIMING_KEY = 'TimingFunction';
var ANIMATION_ITERATION_COUNT_KEY = 'IterationCount';
var ANIMATION_PLAYSTATE_KEY = 'PlayState';
var SAFE_FAST_FORWARD_DURATION_VALUE = 9999;

var ANIMATION_DELAY_PROP = ANIMATION_PROP + DELAY_KEY;
var ANIMATION_DURATION_PROP = ANIMATION_PROP + DURATION_KEY;
var TRANSITION_DELAY_PROP = TRANSITION_PROP + DELAY_KEY;
var TRANSITION_DURATION_PROP = TRANSITION_PROP + DURATION_KEY;

var ngMinErr = angular.$$minErr('ng');
function assertArg(arg, name, reason) {
    if (!arg) {
        throw ngMinErr('areq', 'Argument \'{0}\' is {1}', (name || '?'), (reason || ''));
    }
    return arg;
}

function concatWithSpace(a,b) {
    if (!a) return b;
    if (!b) return a;
    return a + ' ' + b;
}

var helpers = {
    blockTransitions: function(node, duration) {
        // we use a negative delay value since it performs blocking
        // yet it doesn't kill any existing transitions running on the
        // same element which makes this safe for class-based animations
        var value = duration ? '-' + duration + 's' : '';
        applyInlineStyle(node, [TRANSITION_DELAY_PROP, value]);
        return [TRANSITION_DELAY_PROP, value];
    }
};
```

statement

statement

declaration

declaration

statement

Most semistructured merge conflicts happen because developers independently edit the same or consecutive lines of the same method

the probability of creating a merge conflict is approximately the same when editing methods, class fields, and modifier lists

most conflicting merge scenarios, and merge conflicts, involve more than two developers

Understanding Semistructured Merge Conflict Characteristics in Open-Source Java Projects

Paola Accioly · Paulo Borba · Guilherme Cavalcanti

Received: date / Accepted: date

Abstract Empirical studies show that merge conflicts frequently occur, impairing developers' productivity, since merging conflicting contributions might be a demanding and tedious task. However, the structure of changes that lead to conflicts has not been studied yet. Understanding the underlying structure of conflicts, and the involved syntactic language elements might shed light on how to better avoid merge conflicts. To this end, in this paper we derive a catalog of conflict patterns expressed in terms of the structure of code changes that lead to merge conflicts. We focus on conflicts reported by a semistructured merge tool that exploits knowledge about the underlying syntax of the artifacts. This way, we avoid analyzing a large number of spurious conflicts often reported by typical line-based merge tools. To assess the occurrence of such patterns in different systems, we conduct an empirical study reproducing 70,047 merges from 123 GitHub Java projects. Our results show that most semistructured merge conflicts in our sample happen because developers independently edit the same or consecutive lines of the same method. However, the probability of creating a merge conflict is approximately the same when editing methods, class fields, and modifier lists. Furthermore, we noticed that most part of conflicting merge scenarios, and merge conflicts, involve more than two developers. Also, copying and pasting pieces of code, or even entire files, across different repositories is a common practice and cause of conflicts. Finally, we discuss how our results reveal the need for new research studies and suggest potential improvements to tools supporting collaborative software development.

More structure does not always improve merge accuracy (43,509 merges from more than 500 projects)

tools do not often differ
(24% of the ~1K scenarios with conflicts)

semistructured merge reports more false positives (9x, 36)

structured merge has more false negatives (8x, 39)

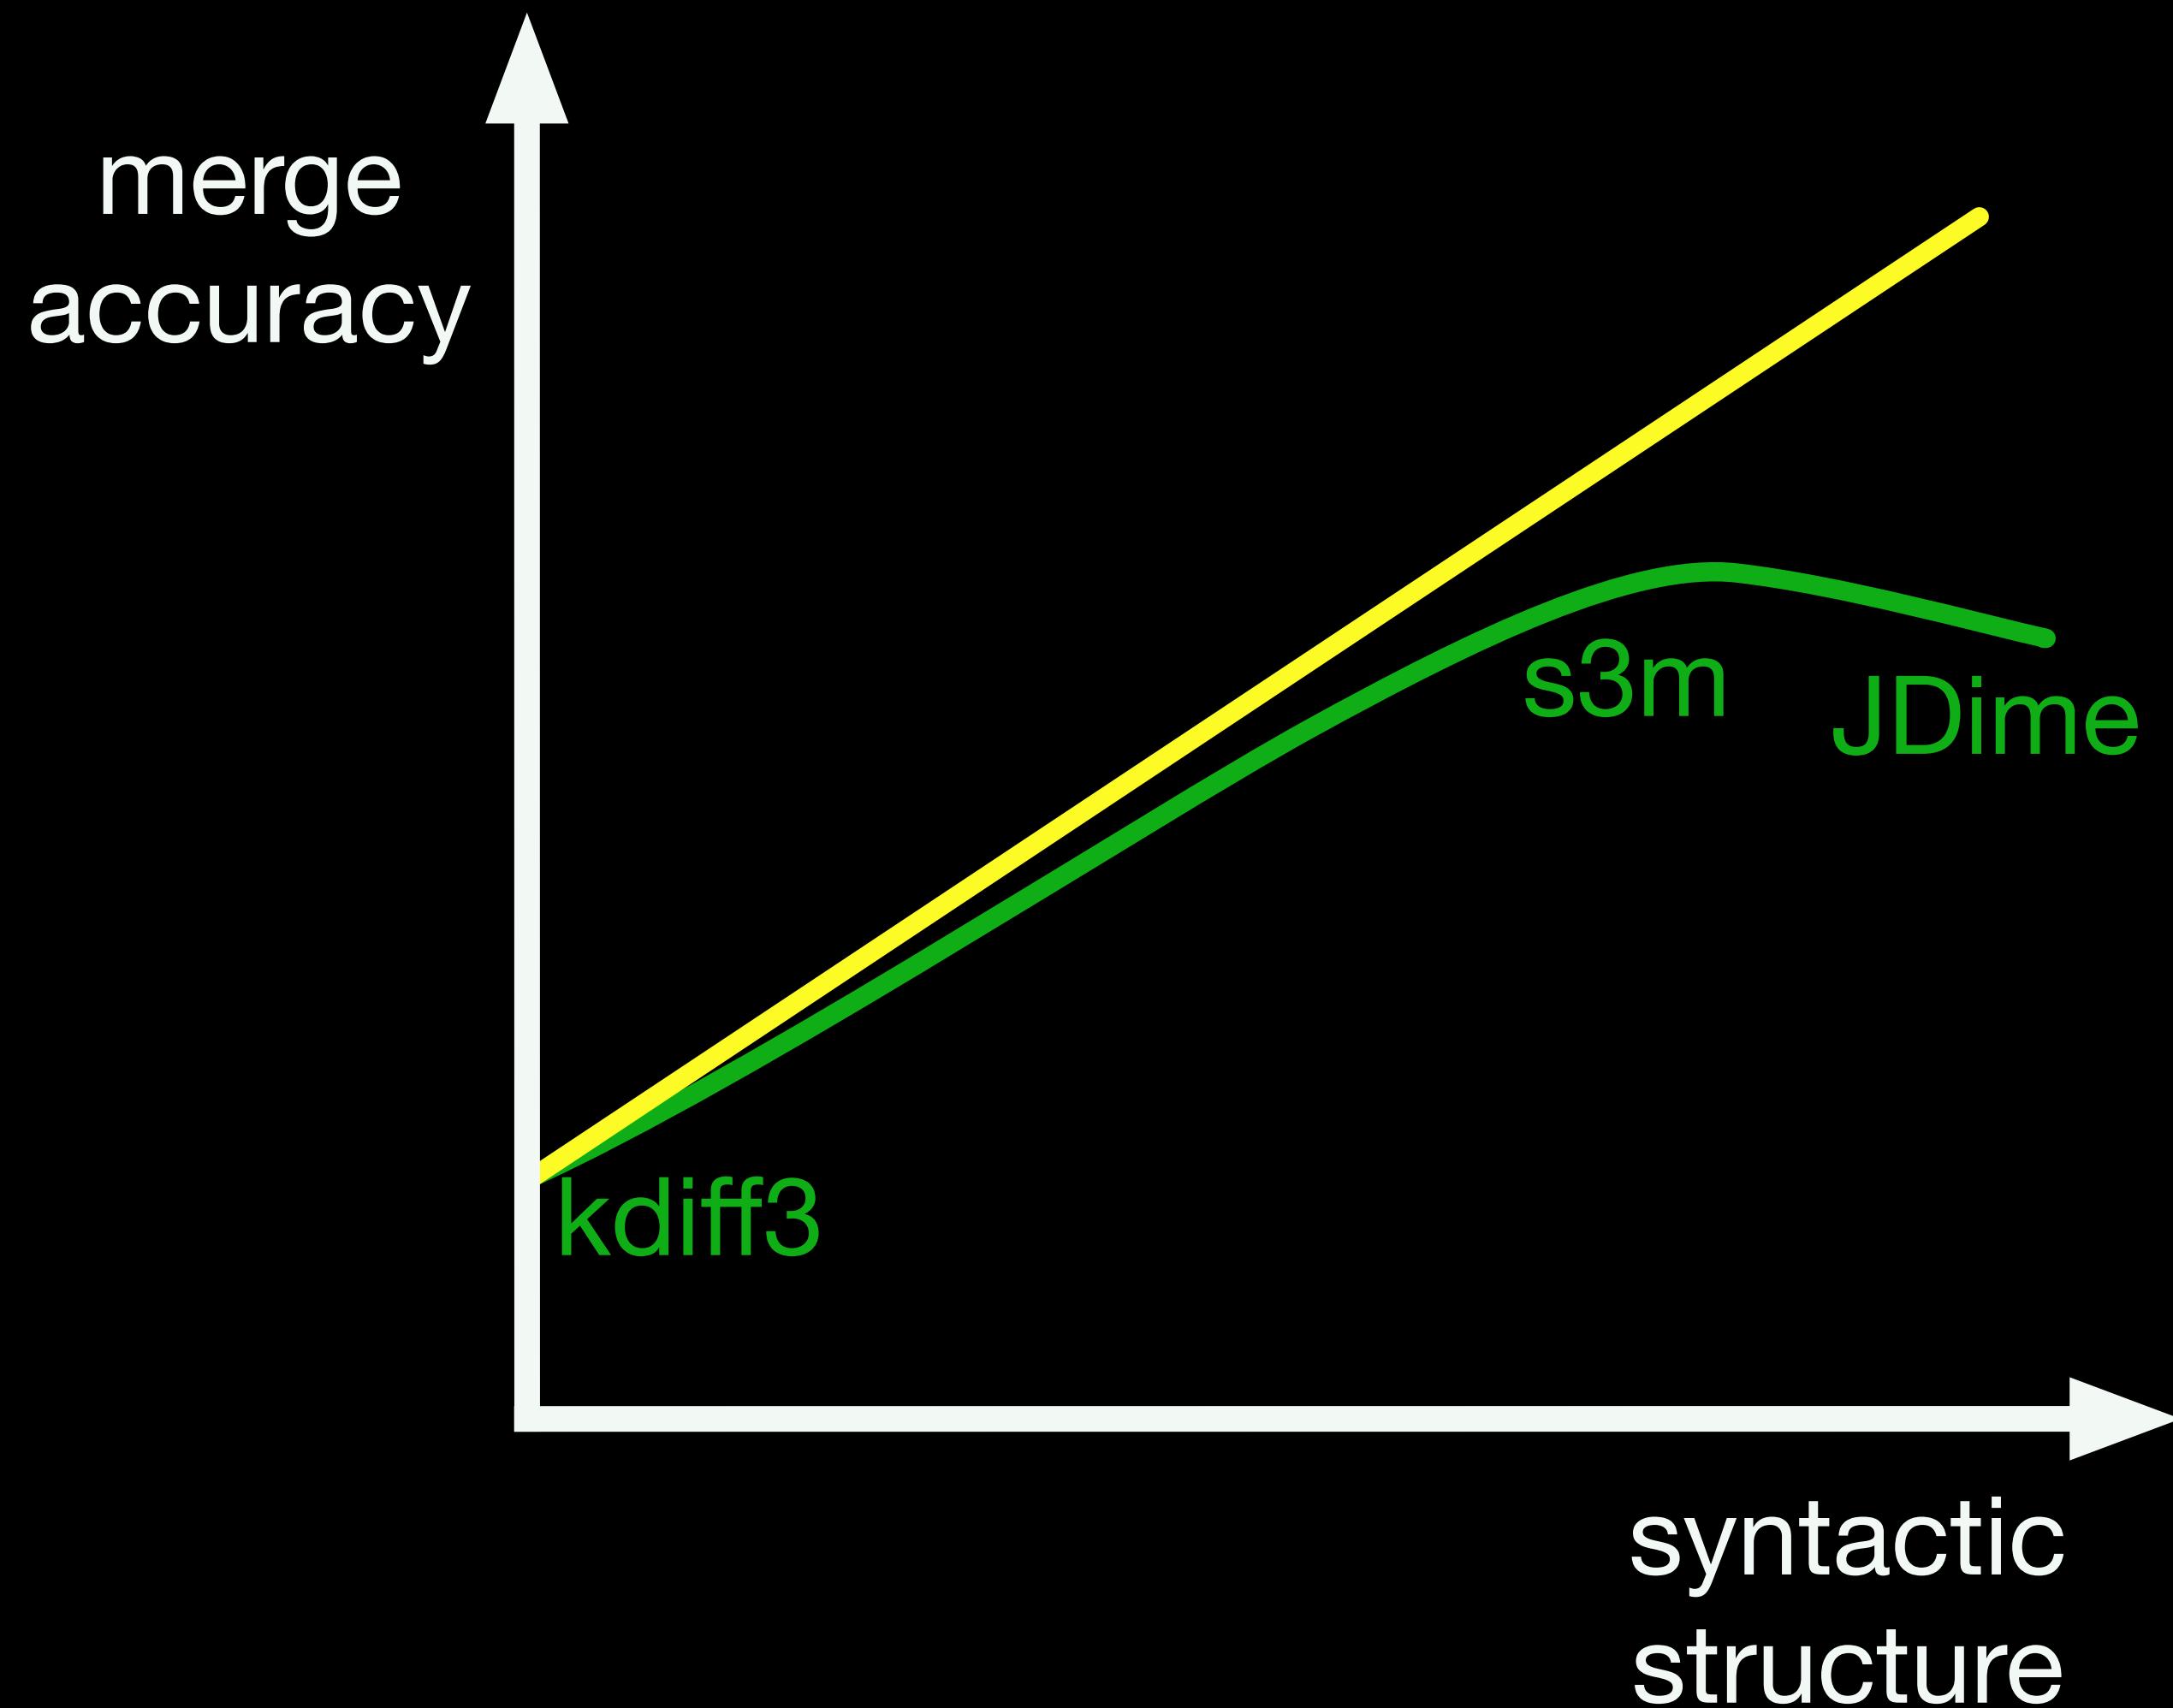
The Impact of Structure on Software Merging:
Semistructured versus Structured Merge

Guilherme Cavalcanti^a, Paulo Borba^a, Georg Seibt^b and Sven Apel^c
^aFederal University of Pernambuco, Recife, Brazil
^bUniversity of Passau, Passau, Germany
^cSaarland University, Saarbrücken, Germany

Abstract—Merge conflicts often occur when developers change the same code artifacts. While state-of-practice unstructured merge tools try to automatically resolve merge conflicts based on textual similarity, semistructured and structured merge tools try to go further by exploiting the syntactic structure of the code. It is not clear whether there is evidence that semistructured merge has significant advantages over unstructured merge, and that structured merge is unknown how semistructured merge compares with structured merge. To help developers decide which kind of tool to use, we conducted a large-scale empirical study. We performed a study by reproducing more than 40,000 merge scenarios from more than 500 projects. In particular, we assess if semistructured merge tools report more conflicts than structured merge tools, if conflicts correctly reported by one but not by the other (false positives), and if conflicts correctly reported by one but missed by the other (false negatives). Our results show that semistructured and structured merge differ in 24% of the scenarios with conflicts. Semistructured merge has more false positives, while structured merge has more false negatives. Finally, we found that adapting a semistructured merge tool to resolve a particular kind of conflict makes semistructured and structured merge even closer.
Index Terms—software merging, collaborative development, code integration, version control systems

G. Cavalcanti, P. Borba, G. Seibt, S. Apel. The Impact of Structure on Software Merging:
Semistructured versus Structured Merge. ASE 2019 <https://pauloborba.cin.ufpe.br/publication/2019the-impact-of-structure-on-softwaremerging-semistructured-versus-structured-merge/>

Literature versus New results suggestions



Combining merge
strategies and avoiding
consecutive line conflicts
show promising results

Sesame = s3m + csdiff

	diff3 vs. SESAME	diff3 vs. s3m	s3m vs. SESAME
Total	1145	1019	675
(%)*	12.04%	10.72%	7.10%

* In relation to the total number of merged files.

Table 3: How often the merge tools differ.

	diff3	SESAME	s3m
Merge conflicts	2413	1413	1632
Conflicting files	1090	657	832

	diff3	SESAME
aFP	291	26
aFN	0	168

	s3m	SESAME
aFP	146	18
aFN	1	48

	diff3	s3m
aFP	169	31
aFN	4	124

Semistructured Merge with Language-Specific Syntactic Separators

Guilherme Cavalcanti
Federal Institute of Pernambuco
Belo Jardim, Brazil

Paulo Borba
Leonardo dos Anjos
Jonatas Clementino
Federal University of Pernambuco
Recife, Brazil

ABSTRACT

Structured merge tools exploit programming language syntactic structure to enhance merge accuracy by reducing spurious conflicts reported by unstructured tools. By creating and handling full ASTs, structured tools are language-specific and harder to implement. They can also be computationally expensive when merging large files. To reduce these drawbacks, semistructured merge tools work with partial ASTs that use strings to represent lower level syntactic structures such as method bodies, and rely on unstructured tools to merge them. This, however, results in merge accuracy loss. To improve accuracy without compromising semistructured merge benefits, we propose a tool that leverages language-specific syntactic separators to infer structure without parsing. We still resort to an unstructured tool to merge lower level structures, but only after preprocessing the code so that text in between separators such as curly braces appear in separate lines. This way we emulate the capabilities of structured merge tools while avoiding their drawbacks. By comparing our tool with a robust implementation of semistructured merge, we find that our tool substantially reduces the number of spurious conflicts. We also observe significant but less substantial reductions on the overall number of reported conflicts, and of files with conflicts. However, similar to structured tools, our tool lets more merge conflicts go undetected. Our tool shows significant improvements over unstructured tools widely used in practice. Finally we observe that exploiting language-specific syntactic separators introduces unique textual alignment challenges.

ACM Reference Format:
Guilherme Cavalcanti, Paulo Borba, Leonardo dos Anjos, and Jonatas Clementino. 2024. Semistructured Merge with Language-Specific Syntactic Separators. In *39th IEEE/ACM International Conference on Automated Software Engineering (ASE '24), October 27–November 1, 2024, Sacramento, CA, USA*. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3691620.3695483>

1 INTRODUCTION

Unstructured merge tools have been in use for over 40 years [1], and are essential for team software development. These tools rely solely on textual analysis to detect and resolve merge conflicts [2]. They are fast and can handle any type of textual content, including source code in any language. However, they often report spurious merge conflicts, and fail to report conflicts that could lead to compilation and execution issues. For example, unstructured merge reports a conflict when asked to integrate rearrangeable declarations (such as method and field declarations) separately added by two developers to the same area of the source code [3]. Simply juxtaposing them in any order would be a correct merge, but an unstructured tool is not able to do that as it recognizes program lines instead of declarations.

To improve merge accuracy, researchers have proposed tools that consider the structure of the source code being merged [3–13]. Such structured merge tools go beyond simple textual, line-based, analysis by parsing the code and handling abstract syntax trees (ASTs). They employ tree matching and combination algorithms. Because of that, they are language-specific and demand significant

Helping people focus
on the merge
conflicts that really
matter!



<http://is.gd/lSOeN7>

Semistructured merge

Merge and Code Review

Paulo Borba
Informatics Center
Federal University of Pernambuco

pauloborba.cin.ufpe.br